

Homework 4

Large-Scale Data Analysis

(individual submission)

Fabian Gieseke & Christian Igel
Department of Computer Science
University of Copenhagen

28 May 2019

Submission Deadline:
11 June 2019, 23:59

This assignment is an **individual assignment**, i.e., you are supposed to solve the assignment on your own. In particular, the report and the code must be written completely by yourself!

Virtual Machine: For the individual assignments, we will generally assume that you make use of the LSDA virtual machine, i.e., all instructions provided are tailored to this system. Of course you are free to use your own desktop environment, tools, and Python installation to implement and test your code. Keep in mind, however, that the output might differ depending on the Python version and the installed packages.

Important: You have to make sure that your code runs on the virtual machine!

Deliverables: Please submit *two separate files* via Absalon: (1) a pdf file `answers.pdf` containing your answers and (2) a `code.zip` file containing the associated code as zip file (do

1 Analyzing Airline Data (50 pts)

We will analyze a dataset containing information about flights in the USA that stems from the *Bureau of Transportation Statistics*.¹ In the zip file provided, you will find a directory `airline_data` containing the files `2016_1.csv`, ..., `2016_6.csv`. Each line of a file (except the headers) contains the flight date, the airline ID, the flight number, the origin airport, the destination airport, the departure time, the departure delay in minutes, the arrival time, the arrival delay in minutes, the time in air in minutes, and the distance between both airports in miles. Make use of your virtual machine and the Hadoop cluster that is already installed on it. For both subtasks, keep in mind to first activate your Python3 environment via `source ~/.venvs/llda/bin/activate`.

1. *Data Analysis with Hadoop:* You first need to start Hadoop by executing `./start_dfs_yarn.sh` from your home directory. Note that both the `mapper.py` and the `reducer.py` need to be executable, which you can achieve via `chmod 755 mapper.py` and `chmod 755 reducer.py`. **Also make sure that the first line of your `mapper.py`/`reducer.py` files equals `#!/usr/bin/env python3`.**

Hint: You can quickly test the mappers and reducers via Unix Streaming, e.g., via:

```
cat 2016*.csv | ./mapper.py | sort | ./reducer.py
```

- (a) *Shortest Flight Distance (5 pts):* Write a mapper and reducer that compute the smallest flight distance per airline ID. The output generated by the reducer should contain, per

¹https://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236&DB_Short_Name=On-Time

line, an airline ID along with the associated minimal flight distance. What is minimal distance for airline ID 20436?

- (b) *Late Arrival Counts (5 pts)*: Write a mapper and reducer that compute the total number of delayed flights per airline ID. Here, a delayed flight is one with a positive arrival delay (treat missing values as 0). The reducer should output, for each airline ID, the number of total flights, the number of delayed flights, and the percentage of delayed flights. What is the percentage of delayed flights for airline ID 20436?
- (c) *Mean and Standard Deviation for Arrival Delays (10 pts)*: Write a mapper and reducer that compute the mean $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$ and the standard deviation $s = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$ for the arrival delays x_1, \dots, x_n for each airline ID. Consider all delay values including the negative ones (treat missing values as 0). A reducer should output, for each airline ID, the associated mean and standard deviation. What is the mean and standard deviation for airline ID 20436?
- (d) *Top-10 of Arrival Delays (10 pts)*: Write a mapper and reducer that compute the 10 most delayed flights for each airline ID. The reducer should output, for each airline ID, a list containing the delays for the 10 most delayed flights. What are the 10 most delayed flights for airline ID 20436?

Note: Write efficient mappers/reducers in the sense that only a constant amount of additional main memory is used by each of the workers. For instance, you can define a couple of “counters” (like for the word count example) or lists of constant size. An example for an inefficient reducer would be a reducer that simply collects all the departure delays per airport in a list to subsequently compute the total departure delay (this list could become very large depending on the amount of incoming data!). In doubt, feel free to get in touch with us!

2. *Data Analysis with Spark*: Implement the four jobs for the Airline dataset in Apache Spark! As a starting point, make use of the Jupyter notebook `Airline Stats.ipynb`, which already creates an RDD based on the airline dataset. You are supposed to implement the tasks in a “similar” fashion as before, i.e., you are supposed to (only) make use of `map` and `reduceByKey` (do not simply make use of some Spark functions that yield, e.g., the mean and standard deviation). In doubt, feel free to ask via Absalon ;-).

Note: You have to execute `source .activate_jupyter_pyspark` before running `pyspark`.

Deliverables

Include answers to the questions raised above to your write-up (`answers.pdf`). Provide, for each subtask, the mapper.py/reducer.py files in a subdirectory (e.g., `1a_shortest_flight_distance`) of your overall source code archive. **Your are also supposed to run all jobs on the Hadoop cluster.** Please add the text output produced by each Hadoop job (e.g., “copy and paste” from the terminal) as well as the generated output files (e.g., `part-00000`) to your `code.zip`. For the Spark task, please add the modified `Airline Stats.HW4.ipynb` notebook to your code.

2 Distributed Linear Regression with Apache Spark (25 pts)

We have briefly discussed how to conduct gradient descent for fitting a logistic regression model in a distributed fashion via Spark. You can find the corresponding notebook on Absalon (`Logistic Regression Spark.ipynb`). This part of the exercise is about implementing linear least-squares regression in a similar fashion!

1. Warm-Up: Follow the code given in the `Logistic Regression Spark.ipynb` notebook.
2. Next, have a look at the first part of the `Distributed_Linear_Least_Squares.ipynb` notebook. The class `LinearLeastSquaresRegression` computes a standard linear regression model (no distributed computations) and resorts to the “black box” optimizer `fmin_l_bfgs_b`

of the `scipy.optimize` package. Note that it is also possible to only provide a function that computes the function values (i.e., the gradient is not necessarily needed). Thus, it is enough to provide a function to the optimizer that computes

$$F(\mathbf{w}) = \frac{1}{n}(\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \quad (1)$$

for a given $\mathbf{w} \in \mathbb{R}^d$, see Slide 19 of Lecture 4.

3. Implement this approach such that the linear least-squares models are computed in a distributed fashion via `map` and `reduce`. More precisely, finalize the implementation of `_function` of the class `LinearLeastSquaresRegressionSpark`, which should return, for a given $\mathbf{w} \in \mathbb{R}^d$, the function value $F(\mathbf{w})$. Note that you can assume that the single vector \mathbf{w} is small enough to be handled by the driver node that executes the optimizer. However, your implementation should be able to deal with a very large number n of data points stored in matrix \mathbf{X} and \mathbf{y} , respectively. Briefly describe the individual steps of your distributed function evaluation (5-10 lines). Add the final figure that is generated to your write-up.

3 Distributed Data Analysis with Apache Spark (25 + ?* pts)

This part will be released next week!