




Faculty of Science



Trees and Forests

Machine Learning

Christian Igel
Department of Computer Science



Outline

- 1 Classification and Regression Trees
- 2 Bias-Variance Decomposition
- 3 Random Forests



Outline

① Classification and Regression Trees

② Bias-Variance Decomposition

③ Random Forests

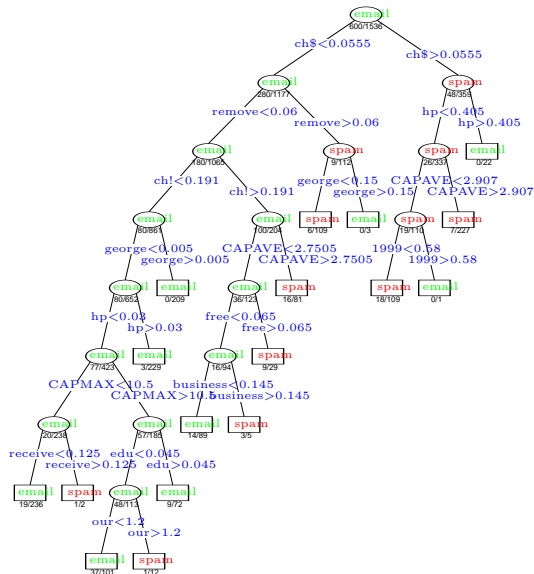


Tree-based methods

- Tree-based models are simple but powerful and are human interpretable.
- Tree-based methods partition the feature space (in the case of the \mathbb{R}^D into rectangular regions) and assign a simple model – usually a constant value/label – to each region \mathcal{R}_τ .
- Several tree-based methods exist (C4.5, ID3, ...), we will focus on CART (Classification and Regression Trees). CART trees are binary trees.



Example: Spam detection



Hastie, Tibshirani, & Friedman. *The Elements of Statistical Learning*. Springer, 2009



Tree evaluation, basic idea

To evaluate a tree given an input x :

- Start at the root node
- Each inner node corresponds to some if-then rule assigning x to one of its children, e.g.:
if $x_d < \theta$ then goto left child node,
else goto the right child node.
- When a leaf node is reached, x is assigned to the value or label (or distribution over labels) associated with that leaf node.

Let the leaf nodes be indexed by $\tau = 1, \dots, |T|$. They define the regions. If x reaches leaf node τ , we have $x \in \mathcal{R}_\tau$.



CART rules

- Every inner node is associated with one coordinate $d \in \{1, \dots, D\}$ and a threshold θ .
- At each inner node, the training data $S = \{(\mathbf{x}_1, y_1), \dots\}$ at that node is split into


$$L_{d,\theta} = \{(\mathbf{x}, y) \in S \mid x_d < \theta\} \text{ and } R_{d,\theta} = \{(\mathbf{x}, y) \in S \mid x_d \geq \theta\} ,$$

passed to the left and right daughter node, respectively.

- The optimal tree cannot be found efficiently. Therefore trees are built using a heuristic.




Example



Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006



Not all splits are possible with a binary tree



Hastie, Tibshirani, & Friedman. *The Elements of Statistical Learning*. Springer, 2009



Building a CART tree, basic idea

- Every inner node is associated with one coordinate $d \in \{1, \dots, D\}$ and a threshold θ .
- At each inner node, the training data $S = \{(\mathbf{x}_1, y_1), \dots\}$ at that node is split into $L_{d,\theta}$ and $R_{d,\theta}$ such that the information gain

$$G_{d,\theta}(S) = Q(S) - \frac{|L_{d,\theta}|}{|S|}Q(L_{d,\theta}) - \frac{|R_{d,\theta}|}{|S|}Q(R_{d,\theta})$$

is maximized, where Q is some impurity measure.

- If the number of datapoints $|S|$ at that node is smaller than a threshold $s_{\text{threshold}}$ or the data is pure (i.e., all elements have the same label), the node becomes a leaf.
- After growing the tree, it is pruned to reduce its complexity.



Simplified objective function

Note that

$$\max_{d,\theta} G_{d,\theta}(S) = Q(S) - \frac{|L_{d,\theta}|}{|S|}Q(L_{d,\theta}) - \frac{|R_{d,\theta}|}{|S|}Q(R_{d,\theta})$$

is equal to

$$\min_{d,\theta} |L_{d,\theta}|Q(L_{d,\theta}) + |R_{d,\theta}|Q(R_{d,\theta})$$



Growing a tree recursively

Procedure GrowTreeRecursively(S)

Input: $S = \{(\mathbf{x}_1, y_1), \dots, \text{maximum number } m \text{ of variables considered per split}\}$

- 1 **if** $|S| < s_{\text{threshold}}$ **then return** *terminal node with labels*
 $\{y_1, \dots, y_{|S|}\}$
 - 2 **if** $\forall (\mathbf{x}_i, y_i), (\mathbf{x}_j, y_j) \in S : y_i = y_j$ **then return** *terminal node with labels* $\{y_1, \dots, y_{|S|}\}$
 - 3 find (d, θ) maximizing $G_{d,\theta}(S)$
 - 4 left child = GrowTreeRecursively ($L_{d,\theta}(S)$)
 - 5 right child = GrowTreeRecursively ($R_{d,\theta}(S)$)
 - 6 **return** *inner node with split* (d, θ)
-



Regression trees

- Training data $\mathcal{S} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\} \in (\mathbb{R}^D \times \mathbb{R})^N$
- Each leaf node $\tau = 1, \dots, |T|$ of a tree T returns a constant, i.e., the output $\hat{y} = T(\mathbf{x})$ given an input \mathbf{x} is

$$T(\mathbf{x}) = \sum_{\tau=1}^{|T|} c_{\tau} \mathbb{I}\{\mathbf{x} \in \mathcal{R}_{\tau}\} \ .$$

- We consider the squared loss $(\hat{y} - y)^2$. Then the choice for the constants minimizing the empirical risk is

$$c_{\tau} = \frac{1}{N_{\tau}} \sum_{\mathbf{x}_n \in \mathcal{R}_{\tau}} y_n$$

with $(\mathbf{x}_n, y_n) \in \mathcal{S}$ and $N_{\tau} = |\{(\mathbf{x}_n, y_n) \in \mathcal{S} \mid \mathbf{x}_n \in \mathcal{R}_{\tau}\}|$.



How to find the optimal split?

- If the impurity measure is the squared loss, we have to find at every node the split (d, θ) solving

$$\min_{d, \theta} \left[\min_{c_L} \sum_{(\mathbf{x}, y) \in L_{d, \theta}} (y - c_L)^2 + \min_{c_R} \sum_{(\mathbf{x}, y) \in R_{d, \theta}} (y - c_R)^2 \right]$$

given the training data at the node.

- The inner minimizations can be solved ... by the averages $c_L = \frac{1}{|L_{d, \theta}|} \sum_{(\mathbf{x}, y) \in L_{d, \theta}} y$ and $c_R = \frac{1}{|R_{d, \theta}|} \sum_{(\mathbf{x}, y) \in R_{d, \theta}} y$.
- For every d , after sorting the training data according to the d th component, only thresholds corresponding to means of d th components of two successive data points need to be tested.



Pruning criterion for regression trees

- For a (sub)tree T , we define the impurity measure

$$Q_{\tau}(T) = \frac{1}{N_{\tau}} \sum_{(\mathbf{x}_n, y_n) \in S \wedge \mathbf{x}_n \in \mathcal{R}_{\tau}} \{y_n - c_{\tau}\}^2, \quad$$

where $N_{\tau} = |\{(\mathbf{x}_n, y_n) \mid (\mathbf{x}_n, y_n) \in S \wedge \mathbf{x}_n \in \mathcal{R}_{\tau}\}|$.

- Pruning criterion for $\alpha \geq 0$:

$$C_{\alpha}(T) = \sum_{\tau=1}^{|T|} N_{\tau} Q_{\tau}(T) + \alpha |T|$$



Pruning regression trees

- **Problem:** For given α , find subtree T_α minimizing $C_\alpha(T)$
- **Solution:** Starting from the full-grown tree T_0 , create sequence of subtrees by collapsing (i.e., replacing by fusing the children) in every step the inner node such that the increase in $\sum_\tau N_\tau Q_\tau(T)$ is minimal
- This finite sequence contains all T_α .
- Proper α is selected using hold-out data.



Classification trees

- Consider classification into K classes
- Let $p_{\tau k}$ be the fraction of training data points in region \mathcal{R}_τ belonging to class k .
- The tree T assigns

$$T(\mathbf{x}) = \operatorname{argmax}_k p_{\tau k} \quad \text{if } \mathbf{x} \in \mathcal{R}_\tau .$$

Note that $p_{\tau k}$ gives a probability distribution over the classes for a given $\mathbf{x} \in \mathcal{R}_\tau$.



Impurity measures for classification trees

Classification error:

$$Q_{\tau}(T) = \frac{1}{N_{\tau}} \sum_{(\mathbf{x}_n, y_n) \in S \wedge \mathbf{x}_n \in \mathcal{R}_{\tau}} \mathbb{I}\{y_n \neq T(\mathbf{x}_n)\}$$

Entropy:


$$Q_{\tau}(T) = - \sum_{k=1}^K p_{\tau k} \ln p_{\tau k}$$

Gini index:

$$Q_{\tau}(T) = \sum_{k=1}^K p_{\tau k} (1 - p_{\tau k})$$




Comparison of impurity measures



binary task, x -axis shows fraction of first class
entropy rescaled to 0.5

Spam classification example: Tree size



Hastie, Tibshirani, & Friedman. *The Elements of Statistical Learning*. Springer, 2009



Outline

- 1 Classification and Regression Trees
- 2 Bias-Variance Decomposition
- 3 Random Forests



Bias-variance decomposition, noiseless case

Let h_S be the hypothesis learnt on training data S and $(x, y) \sim p$. The expected risk under the squared loss is:

$$\begin{aligned}
 & \underbrace{\mathbb{E}_S \mathbb{E}_p \{(y - h_S(x))^2\}}_{\text{risk under squared loss}} \\
 &= \mathbb{E}_S \mathbb{E}_p \{(\underbrace{y - \mathbb{E}_{S'}\{h_{S'}(x)\}}_{\text{bias}} + \underbrace{\mathbb{E}_{S'}\{h_{S'}(x)\} - h_S(x)}_{\text{variance}})^2\} \\
 &= \mathbb{E}_S \mathbb{E}_p \{(\underbrace{y - \mathbb{E}_{S'}\{h_{S'}(x)\}}_{\text{bias}})^2\} \\
 &\quad + 2\mathbb{E}_S \mathbb{E}_p \{(\underbrace{y - \mathbb{E}_{S'}\{h_{S'}(x)\}}_{\text{bias}})(\underbrace{\mathbb{E}_{S'}\{h_{S'}(x)\} - h_S(x)}_{\text{variance}})\} \\
 &\quad + \mathbb{E}_S \mathbb{E}_p \{(\underbrace{\mathbb{E}_{S'}\{h_{S'}(x)\} - h_S(x)}_{\text{variance}})^2\} \\
 &= \underbrace{\mathbb{E}_p \{(y - \mathbb{E}_{S'}\{h_{S'}(x)\})^2\}}_{\text{bias}^2} + \underbrace{\mathbb{E}_S \mathbb{E}_p \{(\mathbb{E}_{S'}\{h_{S'}(x)\} - h_S(x))^2\}}_{\text{variance}}
 \end{aligned}$$

Typically: Large bias indicates that hypothesis class is too restricted, large variance indicates overfitting and occurs if hypothesis class is too complex.



Ensemble methods example

- k regression models
- For each example, model i makes error ϵ_i normally distributed around 0 with variance $\mathbb{E}[\epsilon_i^2] = v$ and covariance $\mathbb{E}[\epsilon_i \epsilon_j] = c$
- We have:

$$\begin{aligned}\mathbb{E} \left[\left(\frac{1}{k} \sum_i \epsilon_i \right)^2 \right] &= \frac{1}{k^2} \mathbb{E} \left[\sum_i \left(\epsilon_i^2 + \sum_{j \neq i} \epsilon_i \epsilon_j \right) \right] \\ &= \frac{1}{k} v + \frac{k-1}{k} c\end{aligned}$$



Outline

- 1 Classification and Regression Trees
- 2 Bias-Variance Decomposition
- 3 Random Forests



Limitations of trees

Single decision trees

- ① are instable, in the sense that changing the data set slightly may lead to a very different tree,
- ② lack smoothness,
- ③ are not well understood in terms of statistical learning theory.

Random forests address the first two issues by averaging over many trees.

- For training, B trees are grown, using different subsets of the data and different splitting variables.
- The outputs of the trees are combined to give the final prediction.

A random forest is an **ensemble classifier**. Averaging models is commonly used for variance reduction.



Growing a random forest tree recursively

Procedure GrowRFTreeRecursively(S, m)

Input: $S = \{(\mathbf{x}_1, y_1), \dots\}$

- 1 **if** $|S| < s_{threshold}$ **then return** *terminal node with labels*
 $\{y_1, \dots, y_{|S|}\}$
 - 2 **if** $\forall (\mathbf{x}_i, y_i), (\mathbf{x}_j, y_j) \in S : y_i = y_j$ **then return** *terminal node*
with labels $\{y_1, \dots, y_{|S|}\}$
 - 3 randomly select variables d_1, \dots, d_m from $\{1, \dots, D\}$ find
 $(d, \theta) \in \{d_1, \dots, d_m\} \times \mathbb{R}$ maximizing $G_{d,\theta}(S)$
 - 4 left child = GrowRFTreeRecursively ($L_{d,\theta}(S)$)
 - 5 right child = GrowRFTreeRecursively ($R_{d,\theta}(S)$)
 - 6 **return** *inner node with split* (d, θ)
-



Growing and evaluating a random forest

Algorithm 1: Random Forest

Input: $S = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$, number of trees B ,
number of variables m

Output: trees T_1, \dots, T_B

```
1 for  $b = 1, \dots, B$  do
2   draw a bootstrap sample  $S'$  by drawing  $N$  elements (with
   replacement) from  $S$ 
3    $T_b = \text{GrowRFTreeRecursively}(S', m)$ 
```

Regression:


$$f_{\text{RF}}(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$$

Classification:

$$f_{\text{RF}}(\mathbf{x}) = \text{majority vote of } T_1, \dots, T_B$$



Example of decision boundary

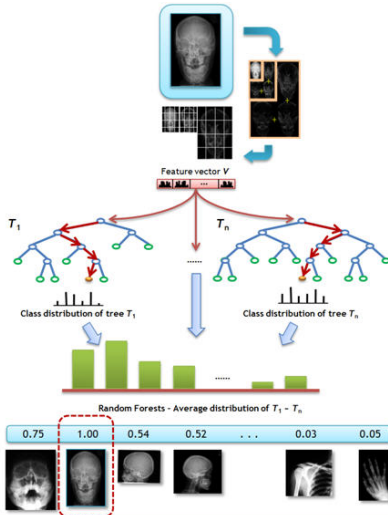


tree (left) vs. random forest (right)

from <http://www.rapidsnail.com>



X-ray classification example



- Number of trees denoted by n instead of B .
- Fusion of the probabilistic output instead of voting. Final results are not normalized.

Ko, Kim, & Nam. X-ray Image Classification using Random Forests with Local Wavelet-Based CS-Local Binary Patterns. *Journal of Digital Imaging* 24(6):1141-1151, 2011



Random forests details

- Pruning is not used.
- Defaults for choosing m :
 - For classification $m = \lfloor \sqrt{D} \rfloor$
 - For regression $m = \lfloor D/3 \rfloor$
- Choosing B : In general, the bigger the better. $B = 100$ may serve as a starting point.




Out-of-bag samples

- Out-of-bag (OOB) samples can be used to evaluate generalization performance of random forests: For each observation x_i in the training set a random forest predictor is built by averaging over all trees constructed not using x_i .
- Number of trees can be determined by increasing forest until OOB sample error converges.
- OOB sample error can be used to determine m instead of using cross-validation.



Business example: Credit scoring


A **credit score** measures the creditworthiness of a client.



figures in this section provided by Kasper Nybo Hansen



Results from MSc thesis



CART: Pros and cons

- ⊕ Good interpretability
- ⊕ Optional probabilistic output
- ⊕ Applicable to both numerical and categorical data
- ⊕ Applicable to large data sets
- ⊖/⊕ Tied to coordinate system
 - ⊖ Suffer from instability and high variance
 - ⊖ Non-smooth decision boundaries
 - ⊖ Comparatively little theoretical understanding in terms of statistical learning theory



Random forests: Pros and cons

- ⊕ Good performance without much tuning
- ⊕ Applicable to both numerical and categorical data
- ⊕ Simple parallelization, applicable to large data sets
- ⊕ Optional probabilistic output
- ⊖/⊕ Tied to coordinate system
- ⊖ Comparatively little theoretical understanding in terms of statistical learning theory

